

Throughput: simulated world-seconds
per wall second

10^2

10^1

2^1

2^3

2^5

2^7

2^9

2^{11}

2^{13}

Number of concurrent worlds

RL training batch
(2048 envs): GPU only

memory wall: 123 MB/world;
1024 worlds need 126 GB (host: 62 GB)

- CENIC reference (Drake, CPU), lockstep batch
- CENIC reference (Drake, CPU), independent worlds
- ◆ CENIC on GPU (Newton/Warp, this work)

